

```
In [15]: import pandas as pd
```

```
df = pd.read_csv(r'C:\Users\humer\Downloads\california.housing.csv
```

```
print(df.head())
print(df.info())
```

```
longitude    latitude    housing_median_age    total_rooms    total_bedrooms
0      -122.23      37.88                  41            880                 1
29.0
1      -122.22      37.86                  21           7099                11
06.0
2      -122.24      37.85                  52           1467                 1
90.0
3      -122.25      37.85                  52           1274                 2
35.0
4      -122.25      37.85                  52           1627                 2
80.0
```

```
population    households    median_income    median_house_value    ocean_proximity
0            322            126        8.3252                  452600
NEAR BAY
1          2401            1138        8.3014                  358500
NEAR BAY
2            496            177        7.2574                  352100
NEAR BAY
3            558            219        5.6431                  341300
NEAR BAY
4            565            259        3.8462                  342200
NEAR BAY
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```


```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	int64
3	total_rooms	20640 non-null	int64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	int64
6	households	20640 non-null	int64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	int64
9	ocean_proximity	20640 non-null	object

```
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB
None
```

```
In [16]: numeric_df = df.select_dtypes(include=['float64', 'int64'])

plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='viridis')
plt.title("Feature Correlation Matrix")
plt.show()
```



```
In [40]: print(df.info())
print(df.describe())

print(df.isnull().sum())

# Correlation heatmap
import seaborn as sns
import matplotlib.pyplot as plt
df_numeric = df.drop(columns=['ocean_proximity'])
sns.heatmap(df_numeric.corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation Heatmap (Numeric Only)")
plt.show()
```

```


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   int64  
 3   total_rooms      20640 non-null   int64  
 4   total_bedrooms   20640 non-null   float64
 5   population       20640 non-null   int64  
 6   households       20640 non-null   int64  
 7   median_income    20640 non-null   float64
 8   median_house_value 20640 non-null   int64  
 9   ocean_proximity  20640 non-null   object  
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB
None
      longitude   latitude  housing_median_age  total_rooms
\ 
count  20640.000000  20640.000000  20640.000000  20640.000000
mean   -119.569704  35.631861   28.639486   2635.763081
std    2.003532    2.135952   12.585558   2181.615252
min   -124.350000  32.540000   1.000000    2.000000
25%   -121.800000  33.930000   18.000000   1447.750000
50%   -118.490000  34.260000   29.000000   2127.000000
75%   -118.010000  37.710000   37.000000   3148.000000
max   -114.310000  41.950000   52.000000   39320.000000

      total_bedrooms  population  households  median_income \
count  20640.000000  20640.000000  20640.000000  20640.000000
mean   536.838857   1425.476744  499.539680   3.870671
std    419.391878   1132.462122  382.329753   1.899822
min   1.000000     3.000000    1.000000   0.499900
25%   297.000000   787.000000   280.000000   2.563400
50%   435.000000   1166.000000  409.000000   3.534800
75%   643.250000   1725.000000  605.000000   4.743250
max   6445.000000  35682.000000  6082.000000  15.000100

      median_house_value
count  20640.000000
mean   206855.816909
std    115395.615874
min   14999.000000
25%   119600.000000
50%   179700.000000
75%   264725.000000
max   500001.000000
longitude          0
latitude           0
housing_median_age 0
total_rooms         0
total_bedrooms     0

```

```
population          0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64
```



In [25]: `from sklearn.linear_model import LinearRegression`

```
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

Out[25]: `LinearRegression`

```
LinearRegression()
```

In [35]: `# Predict using the trained model`

```
y_pred_lr = lr_model.predict(X_test)
```

```
print("Linear Regression Performance:")
print("R2 Score:", r2_score(y_test, y_pred_lr))
print("MAE:", mean_absolute_error(y_test, y_pred_lr))
import numpy as np
```

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred_lr))
print("RMSE:", rmse)
```

Linear Regression Performance:
R² Score: 0.6254240620553608
MAE: 50670.7382409719
RMSE: 70060.52184473517

```
In [34]: from sklearn.ensemble import RandomForestRegressor

rf_model = RandomForestRegressor(random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)

print("Random Forest Performance:")
print("R2 Score:", r2_score(y_test, y_pred_rf))
print("MAE:", mean_absolute_error(y_test, y_pred_rf))
import numpy as np

rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
print("RMSE:", rmse_rf)
```

Random Forest Performance:
R² Score: 0.8169555593071559
MAE: 31643.65566860465
RMSE: 48975.818369986104

```
In [47]: print(X.dtypes)
```

longitude	float64
latitude	float64
housing_median_age	int64
total_rooms	int64
total_bedrooms	float64
population	int64
households	int64
median_income	float64
ocean_proximity	object
dtype:	object

```
In [ ]:
```